Gitansh Wadhwa

gwadhwa@andrew.cmu.edu | +1 (412) 513-9480 | Linkedin://gitanshwadhwa

EDUCATION

Carnegie Mellon University, Pittsburgh, PA — August 2024 - December 2025
Master of Information Systems Management
- Relevant Courses: Cloud Computing, Object Oriented Programming Java, Deep Learning, Distributed Systems
- **Teaching Assistant**: Engineering Data-Intensive Scalable Systems (School of Computer Science, MSE Department)

University of Delhi, New Delhi, IN — August 2019 - July 2023
B.Tech. (Information Technology and Mathematics)
- Relevant Courses: Operating Systems, Data Structures and Algorithm Design, Database Management, Statistics

PROFESSIONAL EXPERIENCE

Amazon, Seattle, WA — May 2025 - Aug 2025
Software Engineer Intern
- **Advertising Infrastructure**: Designed and deployed a low-latency distributed vendor syncing system processing 3Billion+ ad entities daily across Google, Yahoo and Bing with <400ms per entity update
- **Ad Tech Pipelines & Testing:** Built AWS microservices and Scala pipelines with 100% unit test coverage and integration tests to ingest multi-vendor data and resync inconsistencies, reducing sync failures by 85%
- **Campaign Reliability:** Automated anomaly detection workflows, improving advertiser campaign performance by 12%, ensuring consistent ad delivery, and preventing ~$500K in suboptimal ad spend
- **Monitoring & Incident Response:** Implemented CloudWatch alerting, cutting MTTR by 40% and maintaining high-availability for ad services

Binamite, Client Location: Middletown, DE — November 2021 - February 2024
Lead Backend Developer
- **Customer-Facing API Development:** Led the end-to-end development for a fintech application, building RESTful APIs with Spring, leveraging design patterns and Agile practices to deliver a scalable, fully functional MVP
- **Deployment & Reliability:** Built CI/CD pipelines with AWS CodeDeploy and GitHub Actions; integrated automated API and end-to-end tests using Pytest, reducing production bugs by 43% and ensuring safer, faster releases
- **Database Optimization:** Designed and optimized **DynamoDB and Cassandra** tables with efficient partition keys and indexing strategies, reducing query latency by 18% and supporting seamless scalability for high-volume transactions
- **Scalability & Architecture**: Designed and implemented a scalable microservices architecture leveraging AWS, Kafka, Redis and ElasticSearch, enabling handling of over $500K+ in transactions in 2 months of post-beta launch
- **Entrepreneurial Efforts**: Designed and implemented innovative features, conducting technical feasibility analyses and integrating user feedback to align the product with market needs, ensuring a robust and scalable solution

SKILLS

Programming Languages: Python, C++, Java, JavaScript, Typescript, Scala, Bash/Shell, SQL **Web**: Spring, WebSockets, REST APIs, Apache Kafka, Nginx, Spark, HTML/CSS3 **Infrastructure and Databases**: AWS(EC2, Lambda, VPC, S3), CI/CD, Git/Github, Linux/Unix, MySQL, MongoDB, Elastic Search, Redis, Containers (Docker), Terraform, Kubernetes

PROJECTS

AI-Powered Design Generation and Marketplace Integration - Deep Learning (PhD Level) — October 2024
- **Developed a Gen AI platform** that converts user sketches into product-ready images and automates marketplace listings with real-time updates, leveraging custom diffusion models, fine-tuned LLMs to enhance product quality
- **Containerized and deployed a scalable** on a GCP Compute Instance, integrating Stable Diffusion, LLaMA 3, and BART models into a single distributed system, ensuring efficient processing and seamless performance

WeCloud Chat - Microservices Architecture Deployment - (Cloud Computing CMU) — September 2024
- Deployed containerized Login, Chat, and Profile services to Kubernetes clusters on **GCP and AWS**, enabling autoscaling, failover, and fault-tolerant multi-cloud availability
- Built CI/CD pipelines with **Docker, Helm, and GitHub Actions**, streamlining deployments and ensuring consistent multi-cloud delivery